

Data Preparation & Enviroment Setup

```
In [2]: import pandas as pd
import numpy as np
import sqlite3
import matplotlib.pyplot as plt

transactions = pd.read_csv('TRANSACTION TAKEHOME.csv')
product = pd.read_csv('PRODUCTS TAKEHOME.csv')
user = pd.read_csv('USER TAKEHOME.csv')
```

Data Exploration

TRANSACTION TAKEHOME.csv

```
In [4]: # This is the first 10 rows of data, we can get to see the basic sense of the
transactions.head(5)
```

Out[4]:

	RECEIPT_ID	PURCHASE_DATE	SCAN_DATE	STORE_NAME	U
0	0000d256-4041-4a3e-adc4-5623fb6e0c99	2024-08-21	2024-08-21 14:19:06.539Z	WALMART	63b73a7f3d310dcee
1	0001455d-7a92-4a7b-a1d2-c747af1c8fd3	2024-07-20	2024-07-20 09:50:24.206Z	ALDI	62c08877baa38d1a
2	00017e0a-7851-42fb-bfab-0baa96e23586	2024-08-18	2024-08-19 15:38:56.813Z	WALMART	60842f207ac8b7729e
3	000239aa-3478-453d-801e-66a82e39c8af	2024-06-18	2024-06-19 11:03:37.468Z	FOOD LION	63fcd7cea4f8442c33
4	00026b4c-dfe8-49dd-b026-4c2f0fd5c6a1	2024-07-04	2024-07-05 15:56:43.549Z	RANDALLS	6193231ae9b3d7503

```
In [5]: # Display the data structure of the data frame
transactions.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   RECEIPT_ID      50000 non-null  object
1   PURCHASE_DATE   50000 non-null  object
2   SCAN_DATE       50000 non-null  object
3   STORE_NAME      50000 non-null  object
4   USER_ID         50000 non-null  object
5   BARCODE         44238 non-null  float64
6   FINAL_QUANTITY  50000 non-null  object
7   FINAL_SALE      50000 non-null  object
dtypes: float64(1), object(7)
memory usage: 3.1+ MB
```

```
In [6]: transactions['FINAL_QUANTITY'] = transactions['FINAL_QUANTITY'].replace("zero", 0)
transactions['FINAL_QUANTITY'] = pd.to_numeric(transactions['FINAL_QUANTITY'], errors='coerce')
transactions = transactions.dropna(subset=['FINAL_QUANTITY'])

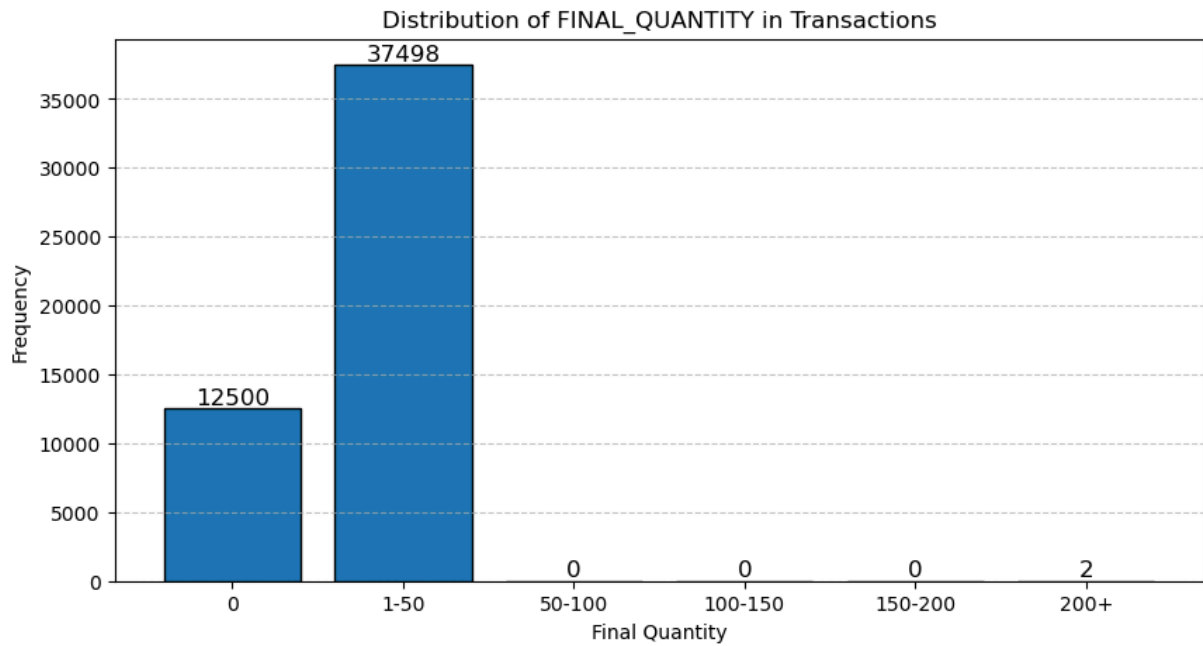
# Plot histogram of FINAL_QUANTITY
bins = [-1, 0, 50, 100, 150, 200, float('inf')]
labels = ['0', '1-50', '50-100', '100-150', '150-200', '200+']
transactions['FINAL_BINNED'] = pd.cut(transactions['FINAL_QUANTITY'], bins=bins, labels=labels)
value_counts = transactions['FINAL_BINNED'].value_counts().sort_index()

plt.figure(figsize=(10, 5))
bars = plt.bar(value_counts.index, value_counts.values, edgecolor='black')

for bar in bars:
    yval = bar.get_height()
    plt.text(bar.get_x() + bar.get_width()/2, yval, int(yval), ha='center', fontweight='bold')

plt.xlabel('Final Quantity')
plt.ylabel('Frequency')
plt.title('Distribution of FINAL_QUANTITY in Transactions')
plt.grid(axis='y', linestyle='--', alpha=0.7)

plt.show()
```



```
In [7]: # Fill the empty space to -1
transactions['FINAL_SALE'] = transactions['FINAL_SALE'].replace(" ", '-1')
# Convert datatype to numeric
transactions['FINAL_SALE'] = pd.to_numeric(transactions['FINAL_SALE'])


# Plot histogram of FINAL_SALE
bins = [-2, -0.5, 0.5, 50, float('inf')]
# Create label for histogram interval
labels = ['Empty', '0', '0', '1-50', '50+']
# Data Visualization
transactions['FINAL_SALE_BINNED'] = pd.cut(transactions['FINAL_SALE'], bins=
value_counts = transactions['FINAL_SALE_BINNED'].value_counts().sort_index()

plt.figure(figsize=(10, 5))
bars = plt.bar(value_counts.index, value_counts.values, edgecolor='black')

for bar in bars:
    yval = bar.get_height()
    plt.text(bar.get_x() + bar.get_width()/2, yval, int(yval), ha='center',

plt.xlabel('Final Sale')
plt.ylabel('Frequency')
plt.title('Distribution of FINAL_SALE in Transactions')
plt.grid(axis='y', linestyle='--', alpha=0.7)

plt.show()
```



Finding:

1. The column of the barcode contains null values.
2. The data type of the "final_quantity" and "final_sale" are objects, which means it contain values other than numeric types.
3. Based on the distribution plot, we find that the "final_sale" contains 0 and is empty when the receipts are scanned. We need to verify the data quality.

Data Cleaning for Transactions Dataset

Assumption: When multiple rows share the same RECEIPT_ID, BARCODE, and USER_ID but have inconsistent values in FINAL_QUANTITY or FINAL_SALE, we assume these discrepancies stem from data entry errors or incomplete records. To ensure data consistency, we convert "zero" in FINAL_QUANTITY to 0 and treat empty FINAL_SALE values as missing data. This approach enables proper aggregation, improves duplicate removal, and ensures we retain the most accurate records. Without external verification, these assumptions allow us to derive the best possible analysis outcomes from the available data.

```
In [10]: # Removal of data that FINAL_SALE have null value
transactions = transactions[transactions['FINAL_SALE'] != -1]
# Removal of data that FINAL_QUANTITY have value of zero
transactions = transactions[transactions['FINAL_QUANTITY'] != 0]
# Check for duplicated datas
duplicate = transactions[transactions.duplicated(keep=False)]
duplicate.shape[0]
```

Out[10]: 276

```
In [11]: # Removal of duplicated rows of data and keep the first data
transactions = transactions.drop_duplicates(keep='first')
# Double check duplicated rows of data
transactions[transactions.duplicated(keep=False)]
```

Out [11]:

RECEIPT_ID	PURCHASE_DATE	SCAN_DATE	STORE_NAME	USER_ID	BARCODE	FI
------------	---------------	-----------	------------	---------	---------	----

In [12]:

```
transactions.drop(columns=['FINAL_BINNED', 'FINAL_SALE_BINNED'], errors='ignore')
#Updated version of transaction dataset. This dataset is well-formatted dataset
transactions.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 24852 entries, 25000 to 49999
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   RECEIPT_ID             24852 non-null  object
1   PURCHASE_DATE          24852 non-null  object
2   SCAN_DATE              24852 non-null  object
3   STORE_NAME             24852 non-null  object
4   USER_ID               24852 non-null  object
5   BARCODE                21996 non-null  float64
6   FINAL_QUANTITY         24852 non-null  float64
7   FINAL_SALE             24852 non-null  float64
dtypes: float64(3), object(5)
memory usage: 1.7+ MB
```

PRODUCT_TAKEHOME.csv

In [14]:

```
# This is the first 5 rows of product data, we can get to see the basic sense of the data
product.head(5)
```

Out [14]:

	CATEGORY_1	CATEGORY_2	CATEGORY_3	CATEGORY_4	MANUFACTURER	BRA
0	Health & Wellness	Sexual Health	Conductivity Gels & Lotions	NaN	NaN	N
1	Snacks	Puffed Snacks	Cheese Curls & Puffs	NaN	NaN	N
2	Health & Wellness	Hair Care	Hair Care Accessories	NaN	PLACEHOLDER MANUFACTURER	ELECS
3	Health & Wellness	Oral Care	Toothpaste	NaN	COLGATE-PALMOLIVE	COLGA
4	Health & Wellness	Medicines & Treatments	Essential Oils	NaN	MAPLE HOLISTICS AND HONEYDEW PRODUCTS INTERCHA...	MAF HOLISTI

In [15]:

```
# Display the data structure of the dataframe
product.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 845552 entries, 0 to 845551
Data columns (total 7 columns):
#   Column          Non-Null Count  Dtype
---  -
0   CATEGORY_1      845441 non-null object
1   CATEGORY_2      844128 non-null object
2   CATEGORY_3      784986 non-null object
3   CATEGORY_4      67459 non-null  object
4   MANUFACTURER    619078 non-null object
5   BRAND           619080 non-null object
6   BARCODE         841527 non-null float64
dtypes: float64(1), object(6)
memory usage: 45.2+ MB
```

Finding:

- Missing value and null value on fields, category / manufacturer / brand / barcode. These missing values can create challenges in accurately linking products to transactions, potentially leading to unmatched records in the transaction data.

```
In [17]: print(f'''CATEGORY_1 contains {product['CATEGORY_1'].dropna().nunique()} distinct categories,
CATEGORY_2 contains {product['CATEGORY_2'].dropna().nunique()} distinct categories,
CATEGORY_3 contains {product['CATEGORY_3'].dropna().nunique()} distinct categories,
CATEGORY_4 contains {product['CATEGORY_4'].dropna().nunique()} distinct categories,
MANUFACTURER contains {product['MANUFACTURER'].dropna().nunique()} distinct manufacturers,
BRAND contains {product['BRAND'].dropna().nunique()} distinct brands ''')
```

```
CATEGORY_1 contains 27 distinct categories,
CATEGORY_2 contains 121 distinct categories,
CATEGORY_3 contains 344 distinct categories,
CATEGORY_4 contains 127 distinct categories,
MANUFACTURER contains 4354 distinct manufacturers,
BRAND contains 8122 distinct brands
```

Data Cleaning for Product Dataset

```
In [19]: # Convert the category with null value to the unknown
product['CATEGORY_1'] = product['CATEGORY_1'].fillna("unknown")
product['CATEGORY_2'] = product['CATEGORY_2'].fillna("unknown")
product['CATEGORY_3'] = product['CATEGORY_3'].fillna("unknown")
product['CATEGORY_4'] = product['CATEGORY_4'].fillna("unknown")
product['MANUFACTURER'] = product['MANUFACTURER'].fillna("unknown")
product['BRAND'] = product['BRAND'].fillna("unknown")
```

```
In [20]: # Double check if the category still contain null values
product.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 845552 entries, 0 to 845551
Data columns (total 7 columns):
#   Column          Non-Null Count  Dtype
---  -
0   CATEGORY_1      845552 non-null object
1   CATEGORY_2      845552 non-null object
2   CATEGORY_3      845552 non-null object
3   CATEGORY_4      845552 non-null object
4   MANUFACTURER    845552 non-null object
5   BRAND           845552 non-null object
6   BARCODE         841527 non-null float64
dtypes: float64(1), object(6)
memory usage: 45.2+ MB
```

USER_TAKEHOME.csv

```
In [22]: # This is the first 10 rows of user data, we can get to see the basic sense
user.head(10)
```

Out [22]:

		ID	CREATED_DATE	BIRTH_DATE	STATE	LANGUAGE	
0	5ef3b4f17053ab141787697d		2020-06-24 20:17:54.000 Z	2000-08-11 00:00:00.000 Z	CA	es-419	
1	5ff220d383fcfc12622b96bc		2021-01-03 19:53:55.000 Z	2001-09-24 04:00:00.000 Z	PA	en	
2	6477950aa55bb77a0e27ee10		2023-05-31 18:42:18.000 Z	1994-10-28 00:00:00.000 Z	FL	es-419	
3	658a306e99b40f103b63ccf8		2023-12-26 01:46:22.000 Z	NaN	NC	en	
4	653cf5d6a225ea102b7ecdc2		2023-10-28 11:51:50.000 Z	1972-03-19 00:00:00.000 Z	PA	en	
5	5fe2b6f3ad416a1265c4ab68		2020-12-23 03:18:11.000 Z	1999-10-27 04:00:00.000 Z	NY	en	
6	651210546816bb4d035b1ead		2023-09-25 22:57:24.000 Z	1983-09-25 22:57:25.000 Z	FL	es-419	
7	642831ea3d4434e63c1936fd		2023-04-01 13:30:18.000 Z	1970-02-16 05:00:00.000 Z	IN	en	
8	63a4c9a1b5f32149b9d82f9e		2022-12-22 21:18:25.000 Z	1982-12-22 05:00:00.000 Z	NC	en	
9	63654b21d02459d8a57a2e2c		2022-11-04 17:25:53.000 Z	1992-05-03 04:00:00.000 Z	NY	en	nor

In [23]:

user['GENDER'].unique()

Out[23]:

array(['female', nan, 'male', 'non_binary', 'transgender',
 'prefer_not_to_say', 'not_listed', 'Non-Binary', 'unknown',
 'not_specified', "My gender isn't listed", 'Prefer not to say'],
 dtype=object)

In [24]:

Display the data structure of the data frame
user.info()


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   ID              100000 non-null object
1   CREATED_DATE    100000 non-null object
2   BIRTH_DATE      96325 non-null  object
3   STATE           95188 non-null  object
4   LANGUAGE        69492 non-null  object
5   GENDER          94108 non-null  object
dtypes: object(6)
memory usage: 4.6+ MB
```

Finding:

- Missing values in the fields, birth date/state/language/gender. For the entry of gender, it contains redundant categories, like "unknown", "not_specified", and "My gender isn't listed" may overlap.

```
In [26]: f'The earliest user creation time is {user['CREATED_DATE'].min()}. The most
```

```
Out[26]: 'The earliest user creation time is 2014-04-18 23:14:55.000 Z. The most rec
ent user creation date is 2024-09-11 17:59:15.000 Z'
```

```
In [27]: f'The oldest user is {user['BIRTH_DATE'].dropna().min()}. The youngest user
```

```
Out[27]: 'The oldest user is 1900-01-01 00:00:00.000 Z. The youngest user creation d
ate is 2022-04-03 07:00:00.000 Z'
```


```
In [28]: gender_counts_all = user['GENDER'].value_counts(dropna=False)

# Replace NaN with "Unknown" for better visualization
gender_counts_all.index = gender_counts_all.index.fillna("unknown")

# Only create label male/female / unknown, as we trying to keep the best eff
labels = [g if g in ["male", "female", "unknown"] else "" for g in gender_cou

plt.figure(figsize=(6, 4))
plt.pie(gender_counts_all, labels=labels, autopct=lambda p: f'{p:.1f}%' if p
plt.title("Distribution of Gender")
plt.show()
```

Distribution of Gender



SQL Queries

1. What are the top 5 brands by receipts scanned among users 21 and over?

Based on the provided transaction records, we could find that COCA-COLA / GREAT VALUE / PEPSI / EQUATE / LAY'S top 5 brands by receipts scanned among users 21 and over. Please find the result under the SQL query.

```
In [30]: conn = sqlite3.connect(":memory:")

user.to_sql("table1", conn, index=False, if_exists="replace")
product.to_sql("table2", conn, index=False, if_exists="replace")
transactions.to_sql("table3", conn, index=False, if_exists="replace")

query = """WITH t1 AS (
    SELECT BARCODE,
           BRAND
    FROM table2),

    t2 AS (
    SELECT ID
    FROM table1
    WHERE (julianday('now') - julianday(BIRTH_DATE)) / 365.25 >= 21),

    t3 AS (
    SELECT t1.BRAND,
           COUNT(table3.RECEIPT_ID)AS number_receipt
    FROM table3
    JOIN t1 ON table3.BARCODE = t1.BARCODE
    LEFT JOIN t2 ON t2.ID = table3.USER_ID
    GROUP BY t1.BRAND)

    SELECT *
```

```

FROM t3
ORDER BY number_receipt DESC LIMIT 6"""

#t1 is the CTE to select the data
#t2 is the CTE to find the users 21 and over
#t3 is the CTE aggregate the number of receipts scanned by the selective use

df_result = pd.read_sql(query, conn)
print(df_result)

```

	BRAND	number_receipt
0	COCA-COLA	535
1	GREAT VALUE	384
2	unknow	366
3	PEPSI	364
4	EQUATE	341
5	LAY'S	324

2. What are the top 5 brands by sales among users that have had their account for at least six months?

Based on the provided transaction records, we could find that PEPSI / COCA-COLA / EQUATE / GREAT VALUE / HERSHEY'S are the top 5 brands by sales among users that have had their account for at least six months. Please find the result under the SQL query.

```

In [32]: query = """WITH t1 AS
              (SELECT DISTINCT ID
               FROM table1
               WHERE (julianday('now') - julianday(CREATED_DATE)) >= 180),

              t2 AS (
                SELECT df2.BRAND,
                       df3.BARCODE,
                       df3.RECEIPT_ID,
                       df3.USER_ID,
                       df3.FINAL_QUANTITY,
                       df3.FINAL_SALE
                FROM table2 df2
                JOIN table3 df3 ON df2.BARCODE = df3.BARCODE
                LEFT JOIN t1 ON df3.USER_ID = t1.ID
              ),

              t3 AS (
                SELECT t2.BRAND,
                       COUNT(t2.RECEIPT_ID) AS num_receipts,
                       SUM(t2.FINAL_QUANTITY*t2.FINAL_SALE) AS total_sales
                FROM t2
                GROUP BY t2.BRAND
                ORDER BY total_sales DESC
                LIMIT 6
              )
              SELECT *
              FROM t3"""

```

#t1 is the CTE to define the account created for at least 6 months

#t2 is the CTE to filter the data in the 6 month period and join the transaction
 #t3 is the CTE to aggregate the sum of sales for each brand

```
df_result = pd.read_sql(query, conn)
print(df_result)
```

	BRAND	num_receipts	total_sales
0	PEPSI	364	3821.3728
1	COCA-COLA	535	3203.6575
2	unknow	366	2449.8067
3	EQUATE	341	2084.1700
4	GREAT VALUE	384	1778.0600
5	HERSHEY'S	190	1430.6771

3. What is the percentage of sales in the Health & Wellness category by generation?

I divided the generation into 6 groups. The oldest fetch user was born in 1900 based on my exploratory data analysis.

Group1(1944 Forward) / Group2(1944-1965) / Group3(1965-1980) / Group4(1980-1995) / Group5(1995-2010) / Group6(2010-Now)

Based on the transaction record, 49.71% of sales come from Group2(1944-1965), 30.09% of sales come from group3(1965-1980), 20.2% of sales come from group4(1980-1995)

```
In [34]: query = """WITH t1 AS
  (SELECT ID,
    birth_year,
    CASE
      WHEN birth_year<1944 THEN 'group1(1944 Forward)'
      WHEN birth_year>=1944 AND birth_year<1965 THEN 'group2(1944-1965)'
      WHEN birth_year>=1965 AND birth_year<1980 THEN 'group3(1965-1980)'
      WHEN birth_year>=1980 AND birth_year<1995 THEN 'group4(1980-1995)'
      WHEN birth_year>=1995 AND birth_year<2010 THEN 'group5(1995-2010)'
      WHEN birth_year>=2010 THEN 'group6(2010-Now)'
      ELSE 'none'
    END AS 'generation'
  FROM(SELECT ID,
    strftime('%Y', BIRTH_DATE)+0 AS birth_year
    FROM table1)
  WHERE birth_year != 'none'),

  t2 AS (
    SELECT t1.generation,
      df3.FINAL_QUANTITY,
      df3.FINAL_SALE
    FROM table2 df2
    JOIN table3 df3 ON df2.BARCODE = df3.BARCODE
    JOIN t1 ON df3.USER_ID = t1.ID
    WHERE df2.CATEGORY_1 = 'Health & Wellness'
  ),

  t3 AS (
    SELECT t2.generation,
      SUM(t2.FINAL_SALE * t2.FINAL_QUANTITY) AS total_sales
```

```

FROM t2
GROUP BY t2.generation
),

t4 AS (
    SELECT SUM(total_sales) AS grand_total
FROM t3
)

SELECT t3.generation,
    t3.total_sales,
    ROUND((t3.total_sales * 100.0 / t4.grand_total), 2) AS percentage
FROM t3
JOIN t4 ON 1=1
ORDER BY t3.total_sales DESC;"""

#t1 is the CTE to define the generation
#t2 is the CTE to filter the data in 'Health & Wellness' and join the user 1
#t3 is the CTE to aggregate the total sales for each generation that have tr
#t4 is the CTE to aggregate the total sales for all generations that have tr
#SQL query to calculate the percentage of sales for each generation

df_result = pd.read_sql(query, conn)
print(df_result)

```

	generation	total_sales	percentage_of_sales
0	group2(1944–1965)	86.56	49.71
1	group3(1965–1980)	52.39	30.09
2	group4(1980–1995)	35.17	20.20

4. Which is the leading brand in the Dips & Salsa category?

Based on the provided transaction records, we could find that TOSTITOS is the leading brand in the Dips & Salsa category which has the most scanned receipts and sales. Please find the result under the SQL query.

```

In [36]: query = """
WITH t1 AS(
    SELECT BARCODE,
        BRAND,
        category_2
    FROM table2
    WHERE category_2 LIKE 'Dips & Salsa'
),

t2 AS(
    SELECT txt1.receipt_id,
        t1.BRAND,
        (txt1.final_quantity*txt1.final_sale)AS sale
    FROM table3 txt1
    JOIN t1 ON txt1.barcode = t1.barcode
    ORDER BY 1,2
),

t3 AS(
    SELECT brand,

```

```

        COUNT(receipt_id) AS receipt_count,
        SUM(sale) AS total_sales
    FROM t2
    GROUP BY brand
)

select *
from t3
order by receipt_count DESC, total_sales DESC
LIMIT 10""""

#t1 is the CTE to find the brand in the Dips & Salsa category
#t2 is the CTE to aggregate the sum of the sales for each brand record in th
#t3 is the CTE to aggregate the number of receipts and the sum of the sales
#Order the t3 receipt_count & total_sales

df_result = pd.read_sql(query, conn)
print(df_result)

```

	brand	receipt_count	total_sales
0	TOSTITOS	36	197.24
1	PACE	24	85.75
2	unknow	21	107.21
3	FRITOS	19	73.76
4	DEAN'S DAIRY DIP	17	39.95
5	MARKETSID	16	65.22
6	HELUVA GOOD!	15	53.98
7	FRESHNESS GUARANTEED	12	46.66
8	MARZETTI	11	51.14
9	HIDDEN VALLEY	10	76.38

Communication with stakeholders

Dear Mr./Mrs,

I hope this message finds you well. I am reaching out to share insights from my exploratory data analysis and to highlight key data quality issues that require further discussion and resolution.

After my exploratory data analysis, there are 3 major data quality issues have been identified.

1. Missing Data

For "PRODUCTS TAKEHOME.csv" has missing values in the category, manufacture, and barcode information.

For "UERS TAKEHOME.csv" has missing values in the information of birth date, state, language, and gender.

For "TRANSACTIONS TAKEHOME.csv" has missing values in the information of barcode.

2. Inconsistent Data Types

For "TRANSACTIONS_TAKEHOME.csv", the column of 'FINAL_QUANTITY' contains non-numeric values like "zero". The column of 'FINAL_SALE' contains many rows of blank space data, making it difficult to determine transaction amounts.

3. Potential of Duplicate Data in "TRANSACTIONS_TAKEHOME.csv"

We found that 50,000 rows in the dataset contain duplicate records when checking the rows: RECEIPT_ID, PURCHASE_DATE, SCAN_DATE, STORE_NAME, USER_ID, and BARCODE.

Based on the SQL query investigation, we found that Coca-Cola and Pepsi are the 2 brands that ranked high in the loyal Fetch user and 21+ years old Fetch users. In other words, these 2 brands have a long-term rewards points campaign event in Fetch Rewards.

In addition to that I would like to address my questions and potential solutions for the outstanding issues,

1. Unknown brand and category classification

Resolution: Can we verify the data and provide the reference data for the unknown products? Before verification, I propose standardizing all unknown brands and categories by overwriting them as "unknown" for consistency.

2. Change of FINAL_SALE & FINAL_QUANTITY for transactions data

Resolution: Can we confirm if "zero" values and blank spaces represent actual data? However, before we verify the data, I would like to overwrite the 'zero' to integer 0 for numeric consistency. Also, I would like to overwrite the " " to integer 0.

3. Duplicated Data

Resolution: Can we verify if duplicate transactions should be merged or if they result from ingestion errors?(check for FINAL_SALE & FINAL_QUANTITY) and merge the duplicate row? Before verification, I would like to drop the data that have 'zero' in 'FINAL_QUANTITY' and " "(empty space) in 'FINAL_SALE'. Since we could not make any assumption based on the unknown situation. Lastly, we need to review duplicate rows that occur more than twice and retain only the first occurrence in the dataset.

I am looking forward to hearing your thoughts and clarifications. Please let me know how you'd like to proceed, and I'd be happy to make any necessary adjustments.

Best Regards,

Xinghong Ma